

www.storagedeveloper.org

OpenSDS Manageability using Swordfish for Cloud-native Frameworks

**Anjaneya 'Reddy' Chagam, Sr. Principal Engineer, Intel Corporation
Steven Tan, VP & CTO Cloud Storage, Huawei**



Notices & Disclaimers



Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Performance varies depending on system configuration. **No product can be absolutely secure.**

Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. For more complete information about performance and benchmark results, visit <http://www.intel.com/benchmarks>.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors. Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. For more complete information visit <http://www.intel.com/benchmarks>.

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

The benchmark results may need to be revised as additional testing is conducted. The results depend on the specific platform configurations and workloads utilized in the testing, and may not be applicable to any particular user's components, computer system or workloads. The results are not necessarily representative of other benchmarks and other benchmark results may show greater or lesser impact from mitigations.

Intel does not control or audit third-party benchmark data or the web sites referenced in this document. You should visit the referenced web site and confirm whether referenced data are accurate.

© 2018 Intel Corporation.

Intel, the Intel logo, and Intel Xeon are trademarks of Intel Corporation in the U.S. and/or other countries.

*Other names and brands may be claimed as property of others.

Agenda

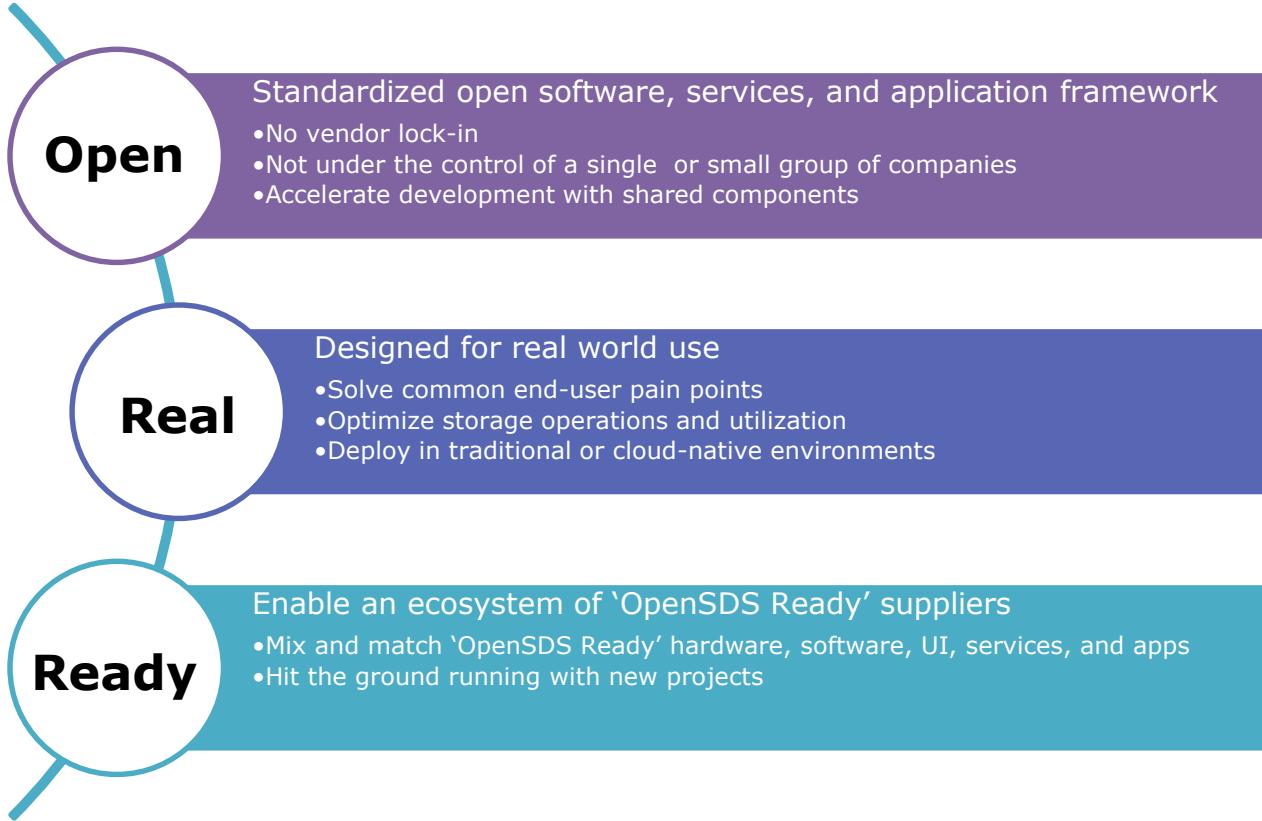
- ❑ OpenSDS Overview & Architecture
- ❑ OpenSDS Kubernetes Integration Architecture
- ❑ OpenSDS Swordfish Manageability
- ❑ OpenSDS Flash Management
- ❑ Summary
- ❑ Q&A

Overview



An open source community working to address **data storage integration** challenges, particularly in **scale-out cloud native environments** with **heterogeneous storage platforms**.





Goals

Open

Standardized open software, services, and application framework

- No vendor lock-in
- Not under the control of a single or small group of companies
- Accelerate development with shared components

Real

Designed for real world use

- Solve common end-user pain points
- Optimize storage operations and utilization
- Deploy in traditional or cloud-native environments

Ready

Enable an ecosystem of 'OpenSDS Ready' suppliers

- Mix and match 'OpenSDS Ready' hardware, software, UI, services, and apps
- Hit the ground running with new projects

Overview

 Open SDS Platform	Control/ Management Plane	Standard REST API, Single Pane Management, Policy-Based, Storage and Data Services, Orchestration and Automation
	Data Plane	Data Reliability, High Availability, Data Protection, Data Mobility, Data Reduction, I/O Performance
	Storage/ Service Layer	DAS (HDD/Flash), SAN, NAS, AFA, Commodity, Private Cloud, Public Cloud

Decoupled control plane from data plane and storage/service layer

Manageability Focus



Orchestration & Automation

Data lifecycle, protection, replication, migration, security, governance, optimization



Multi-Cloud Data Management

Placement, access, security, and search across private and public clouds - AWS, Azure, GCP, and more



Data & Storage Intelligence

Data collection, monitor, reports, analytics, forecast, ML/AI

The Core Projects

SUSHI

The Northbound Plug-ins Project

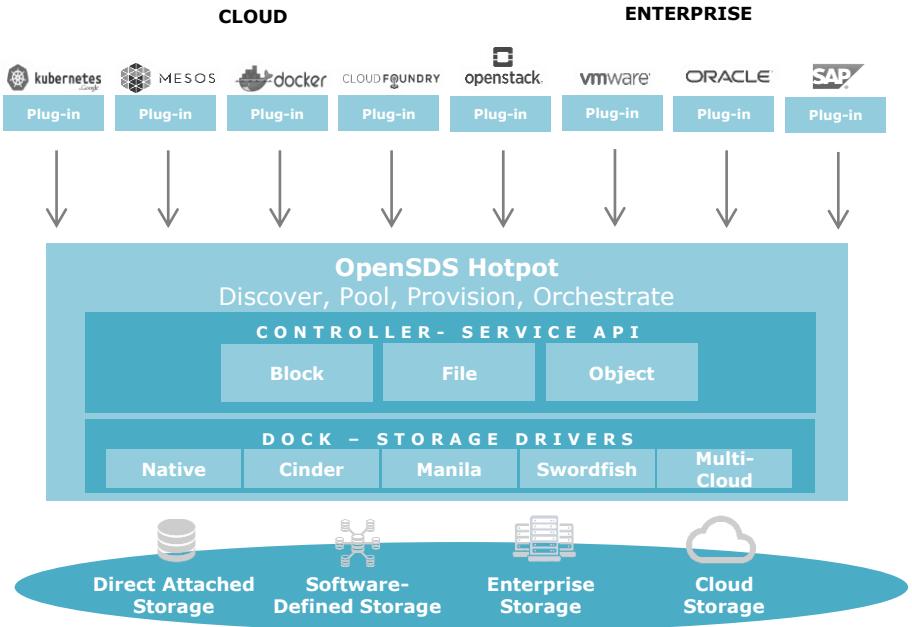
Common plug-ins to enable OpenSDS storage services for cloud and application frameworks



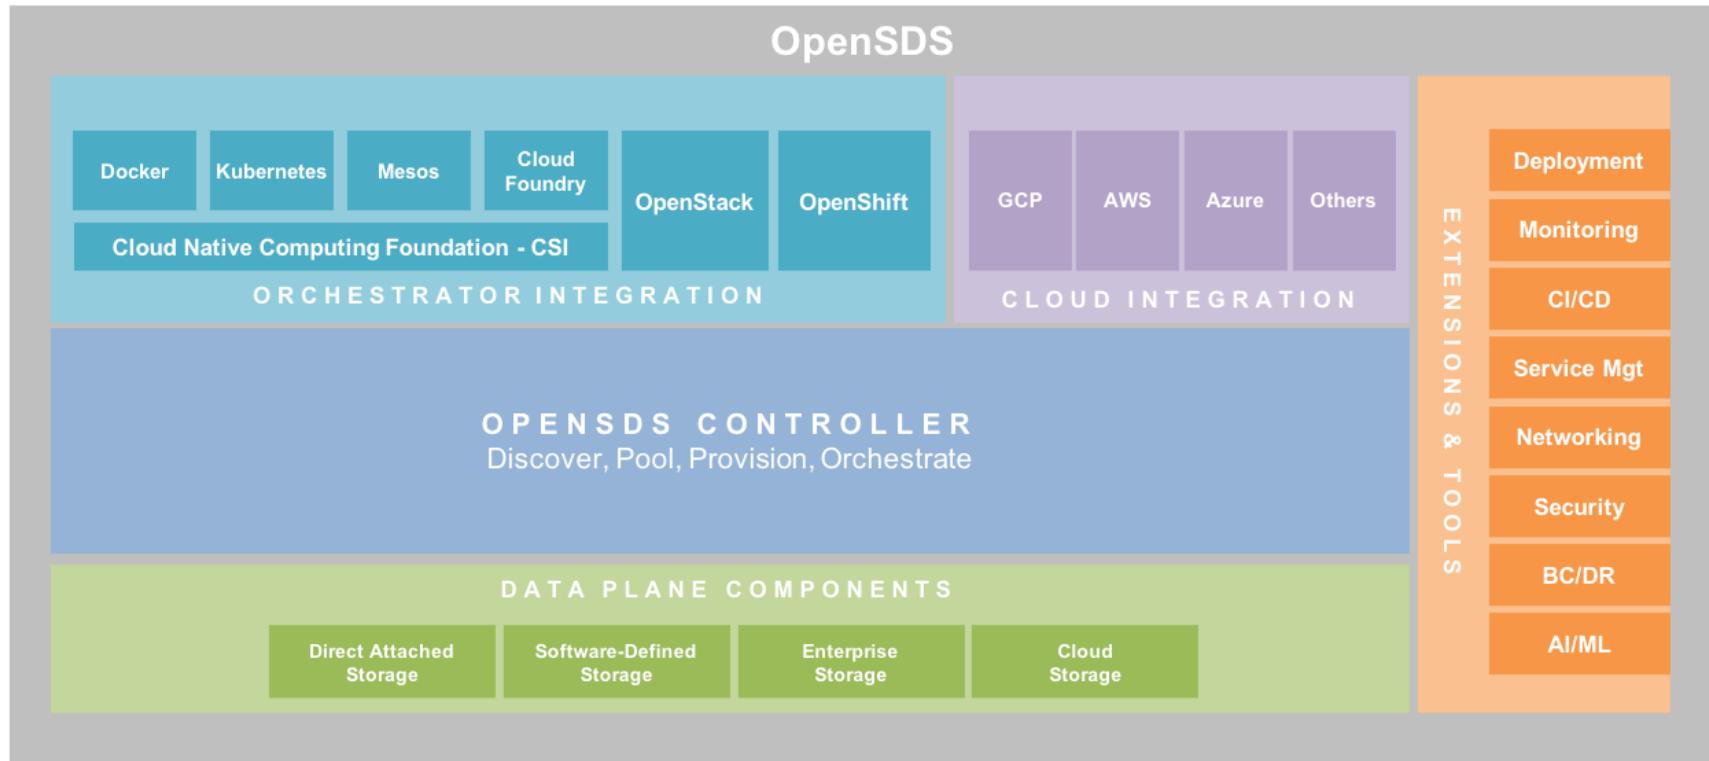
HOTPOT

The Storage Controller Project

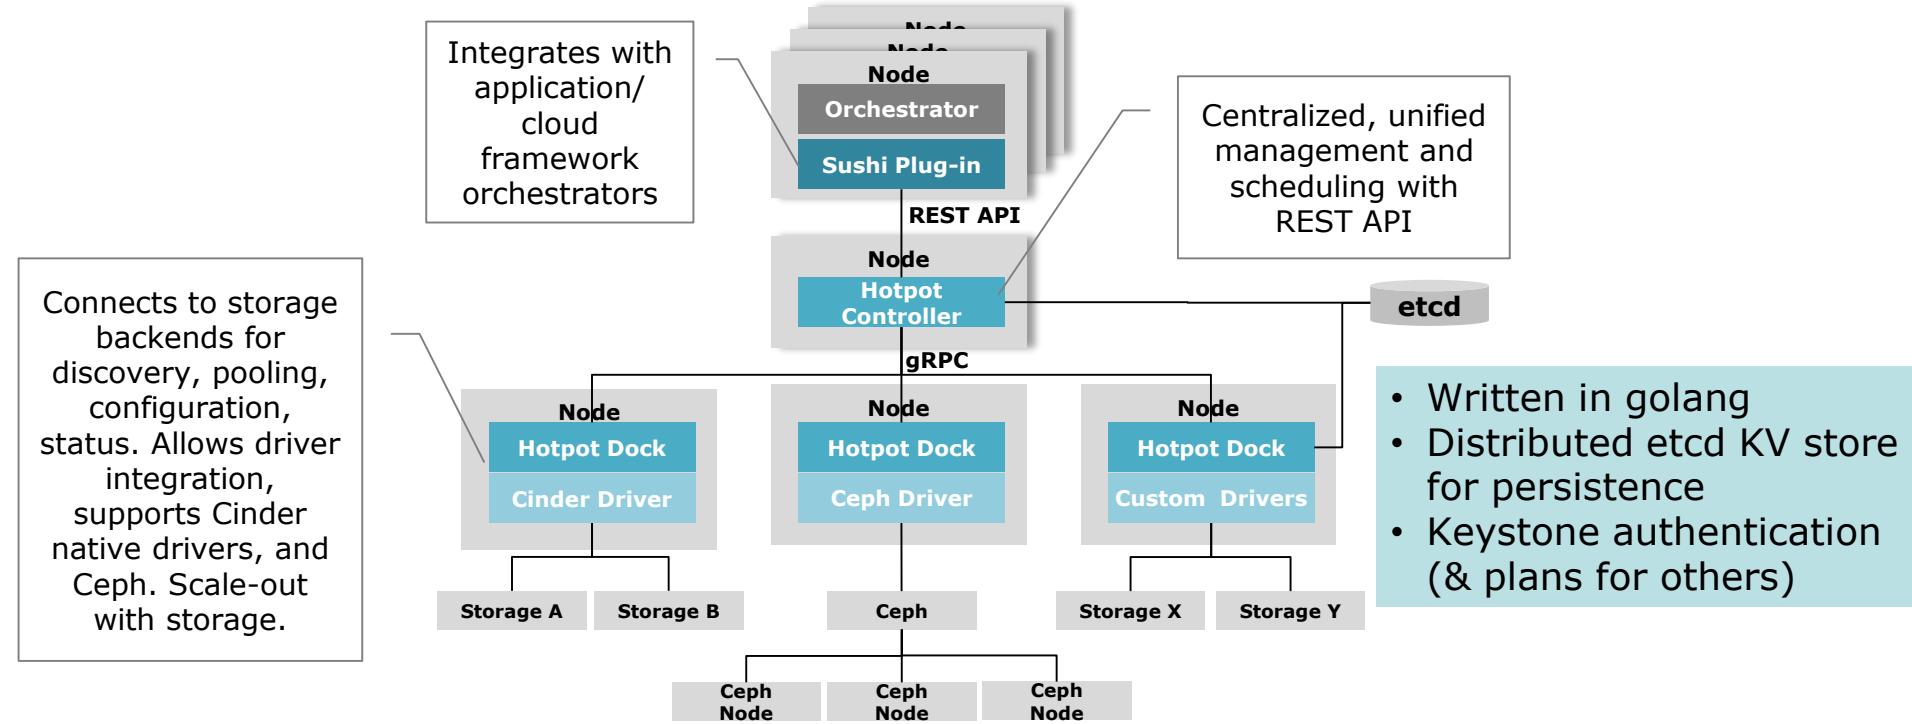
Single control for block, file, and object services across storage on premise and in clouds



OpenSDS Framework



OpenSDS Architecture



The Dashboard

admin
Region_Chengdu

Home
Update 5 minutes ago

Volume
3 volumes

Profile
2 profiles have been created

Resource
2 storages, 2 availability zone

Identity
Managing tenants and users

Create Profile

Create profile to suit your specific use requirements.

* Name: high_performance

* Access Protocol: iSCSI

* Provisioning Type: Thin Auto Thick

QoS Policy: Enable

MaxIOPS	8000	IOPS/TB
MBWS	100	MBWS/TB

Replication Policy: Enable

Snapshot Policy: Enable

Customization: Add

Available Storage Pool: Matching

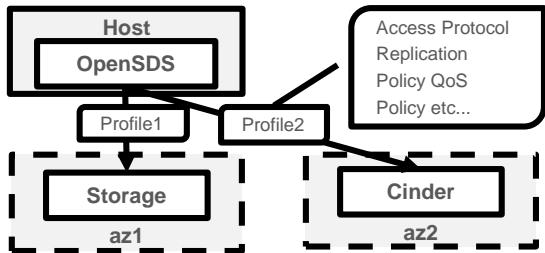
OK Cancel

Sample End User Scenarios

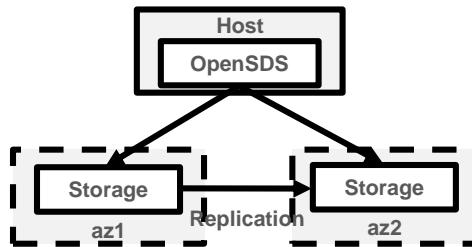
1. **Consolidation** of storage control and management across OpenStack and K8S clusters
2. **Snapshot backups** to cloud and restore to any storage platform
3. **Host-based replication** across different storage backends
4. **Data lifecycle management** to move data across storage tiers and archival for long term retention
5. **Multi-cloud data control** to allow the use of any cloud object store
6. **Orchestrate and automate** provisioning and backups across different storage platforms and backup software

OpenSDS Use Cases

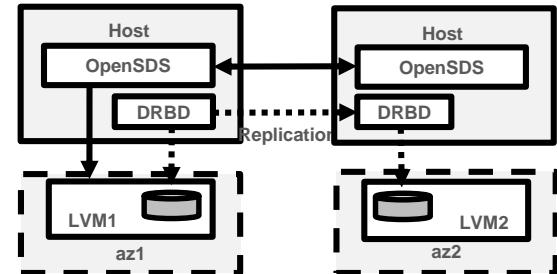
1. Basic Operations(multi storage backend)



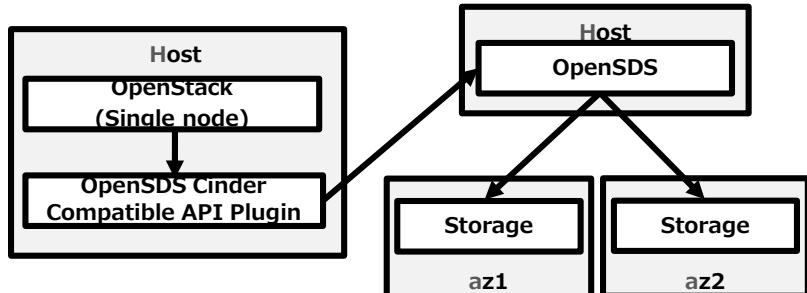
2. Array-based Replication



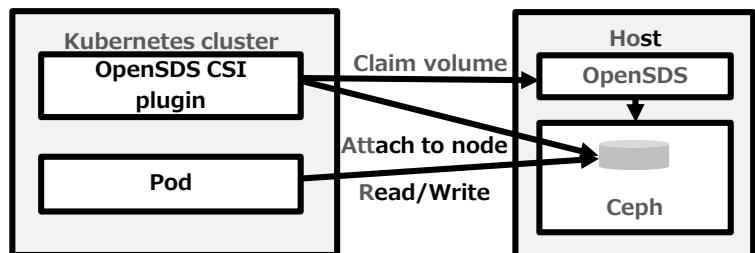
3. Host-based Replication



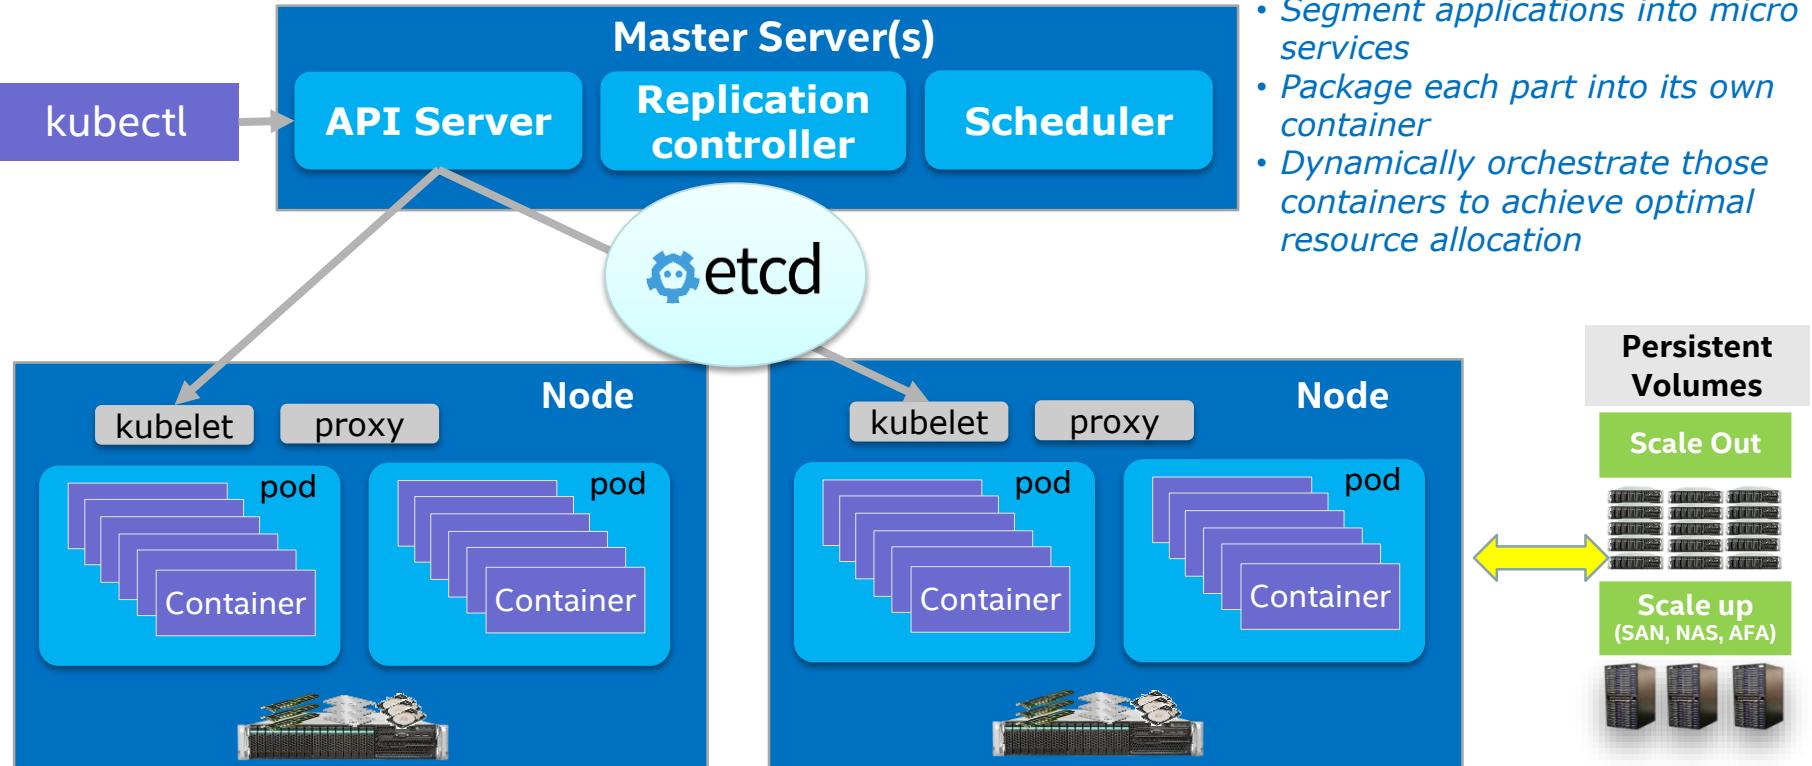
4. OpenStack Integration



5. Kubernetes Integration

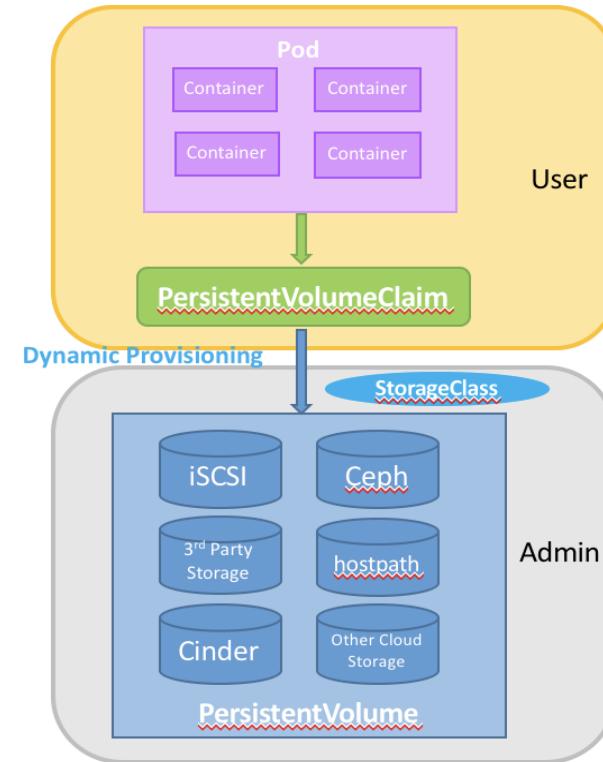


Kubernetes Overview

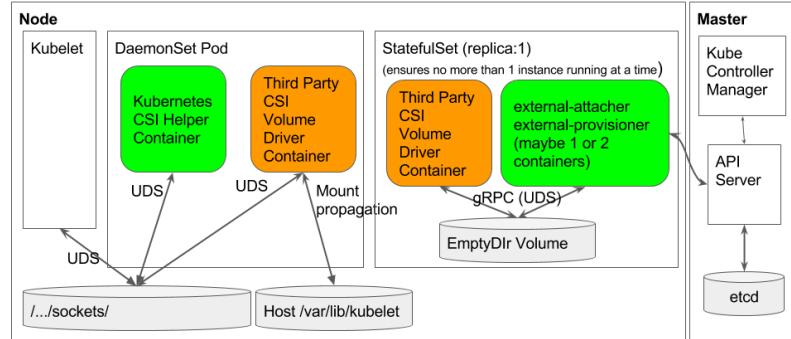
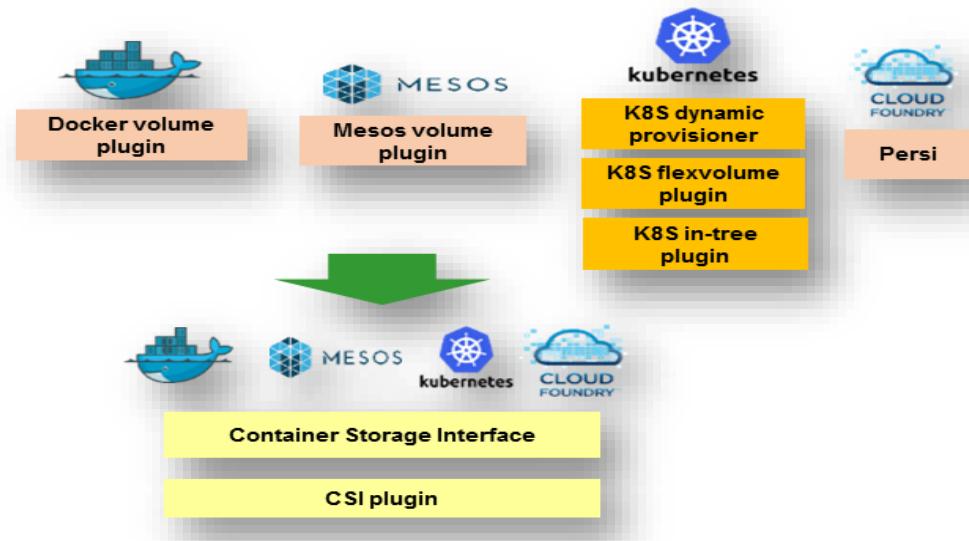


Kubernetes Persistent Volumes

- A **PersistentVolume** (PV) is a piece of storage in the cluster that has been provisioned by an administrator.
- A PV can be provisioned **statically** or **dynamically**.
- A **PersistentVolumeClaim** (PVC) is a request for storage by a user through a StorageClass.
- A **StorageClass** provides a way for administrators to describe the “classes” of storage they offer. Different classes might map to different quality-of-service levels (or “profiles”) in other storage systems.
- A **StorageClass** needs to specify a provisioner for dynamic provisioning.



Container Storage Interface (CSI)



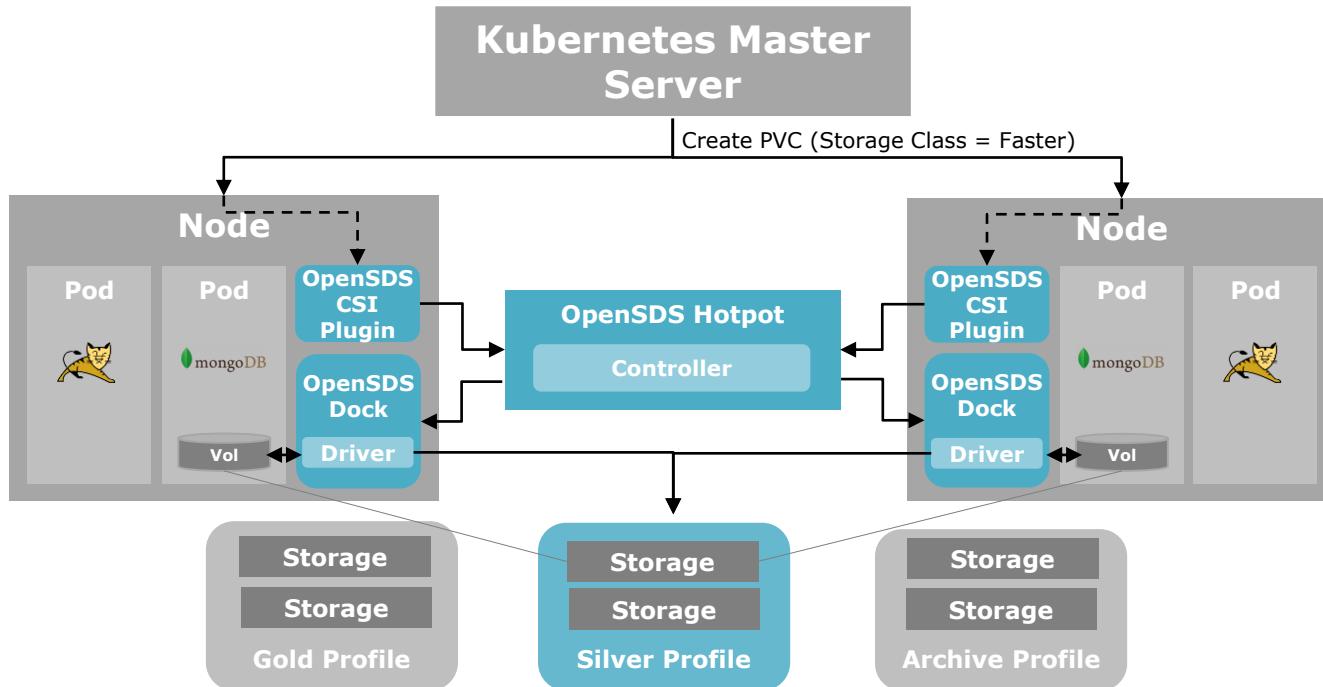
Legend:

- Orange box: External Component - Created by Third Party Storage Vendor
- Green box: External Component - Created by Kubernetes Team

Source: <https://github.com/kubernetes/community/blob/master/contributors/design-proposals/storage/container-storage-interface.md>

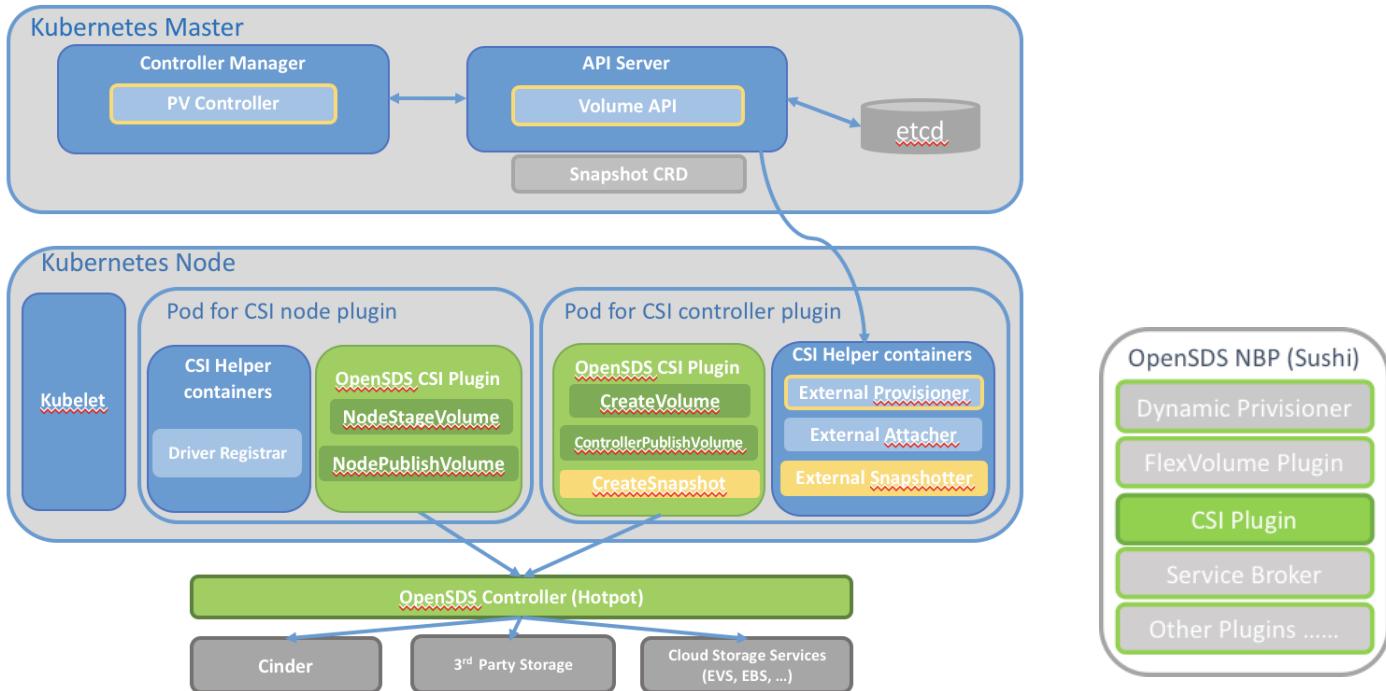
CSI is an industry standard defined to enable storage vendors to develop a plugin once and have it work across a number of container orchestration systems

OpenSDS Kubernetes Architecture

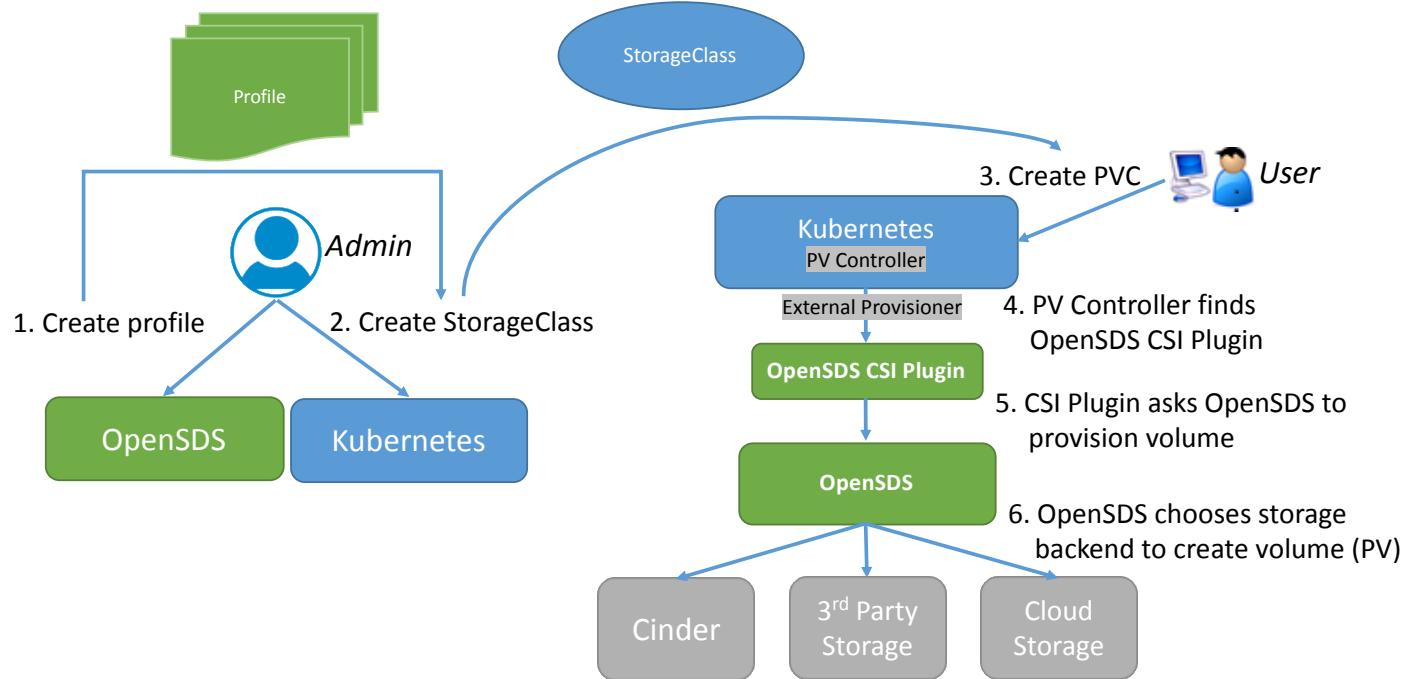


OpenSDS provides driver integration and rich storage services (e.g. data protection, disaster recovery)

OpenSDS - Provision Volumes



OpenSDS – Profile Mapping to K8S

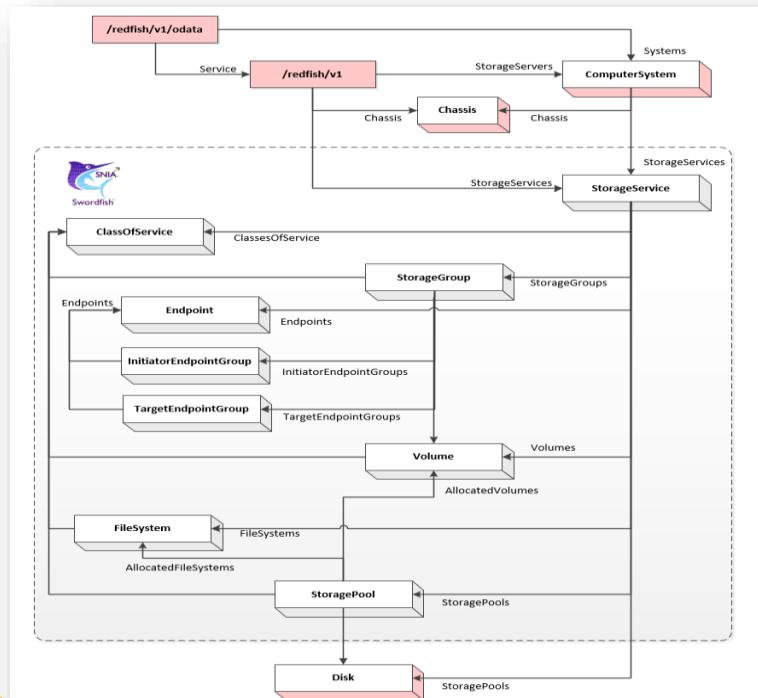


Profile - Policy Driven Management



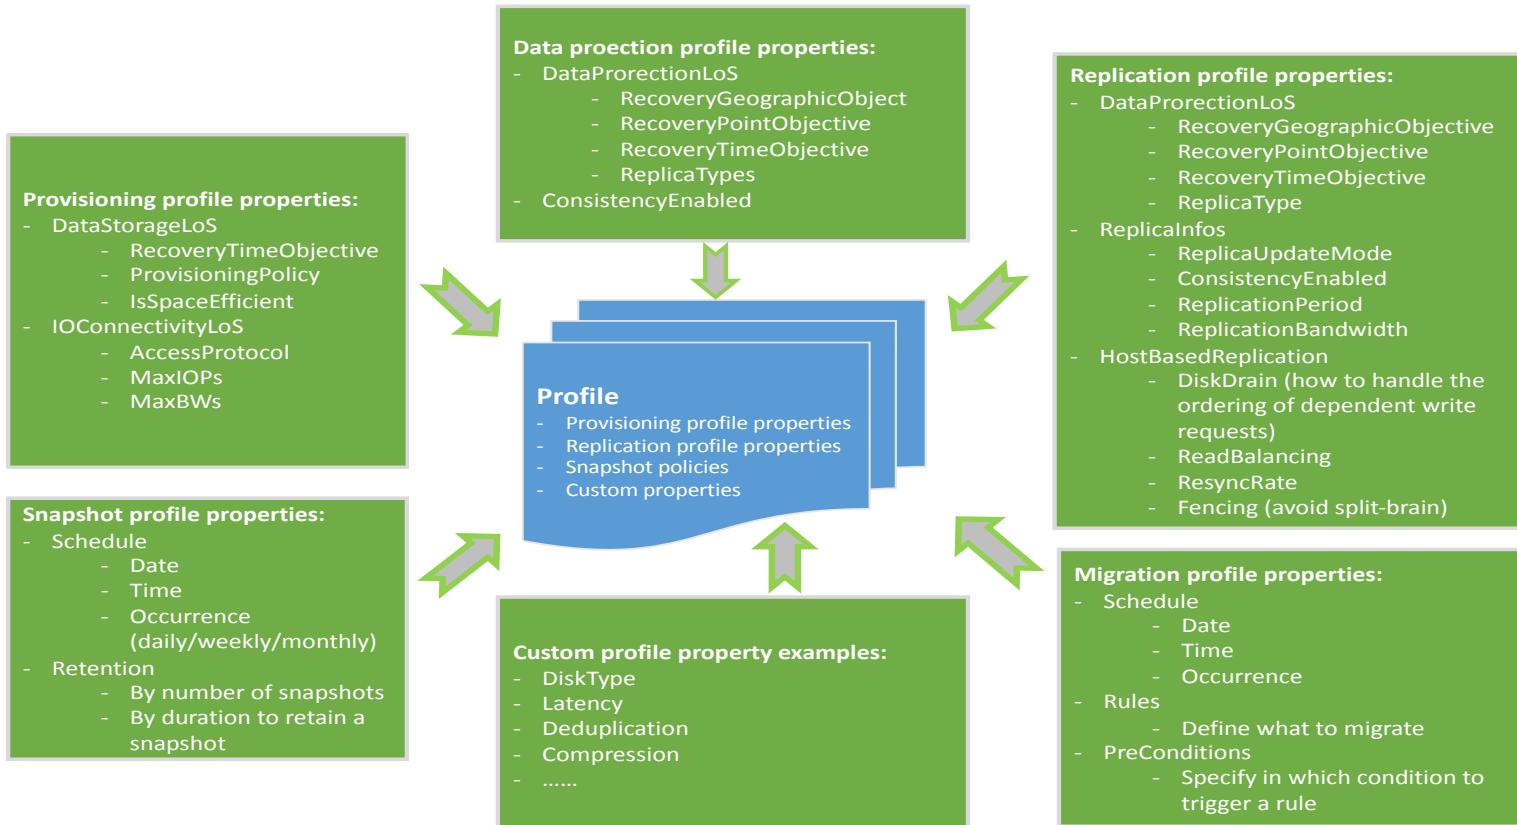
- OpenSDS profile is based on Swordfish specification.
- The SNIA Swordfish™ specification helps to provide a unified approach for the management of storage and servers in hyperscale and cloud infrastructure environments, supported by multiple storage vendors.
- An extension of the DMTF (Distributed Management Task Force) Redfish specification.
 - Redfish is designed by the DMTF's Scalable Platforms Management Forum (SPMF) to create and publish an open industry standard specification and schema for management of scalable platform hardware. It is a RESTful interface over HTTPS in JSON format based on OData v4.

Profile is the key abstraction in “OpenSDS” for storage provisioning to support multi-cloud orchestration



Source: Swordfish_v1.0.5_Specification

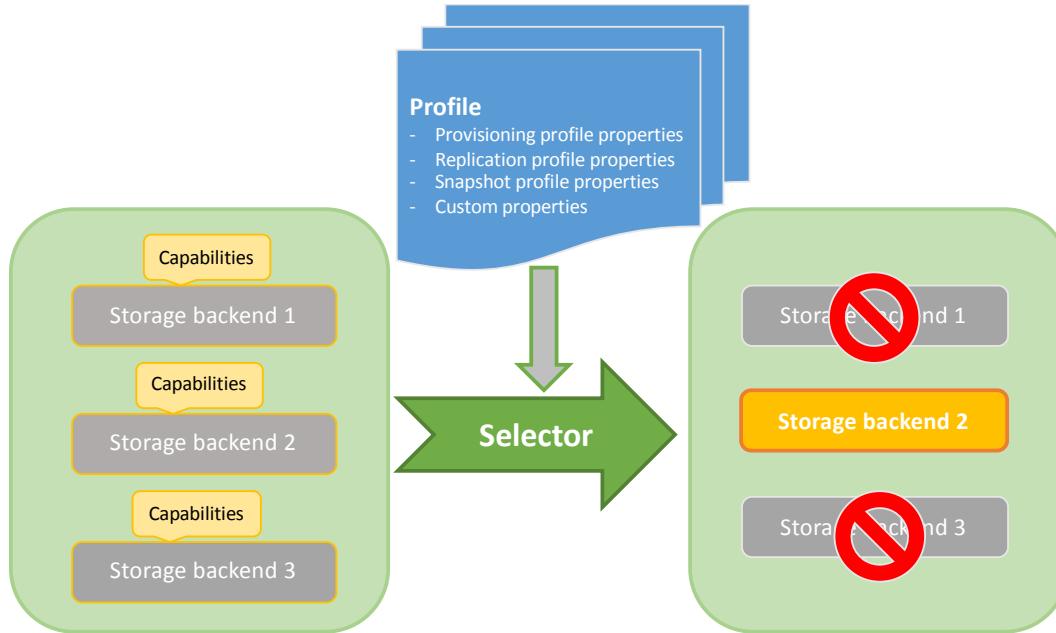
Profile - Swordfish Resource Mapping



Reference: <https://docs.google.com/document/d/1irNnz019j0XuW6SZNigs6QuYFOCC3uL44EkRYJiHyq8/edit#>

2018 Storage Developer Conference. © Intel Corporation. © Huawei. All Rights Reserved.

Mapping Profile to Storage Systems



OpenSDS K8S CSI Plug-in

HighPerformanceSC.yaml

```
apiVersion: storage.k8s.io/v1
```

```
kind: StorageClass
```

```
metadata:
```

```
  name: opensds-csi-high-performance-sc
```

```
provisioner: csi-opensdsplugin
```

```
parameters:
```

```
  profile: High-Performance
```

HighPerformancePVC.yaml

```
apiVersion: v1
```

```
kind: PersistentVolumeClaim
```

```
metadata:
```

```
  name: opensds-csi-high-performance-pvc
```

```
spec:
```

```
  accessModes:
```

```
    - ReadWriteOnce
```

```
  resources:
```

```
    requests:
```

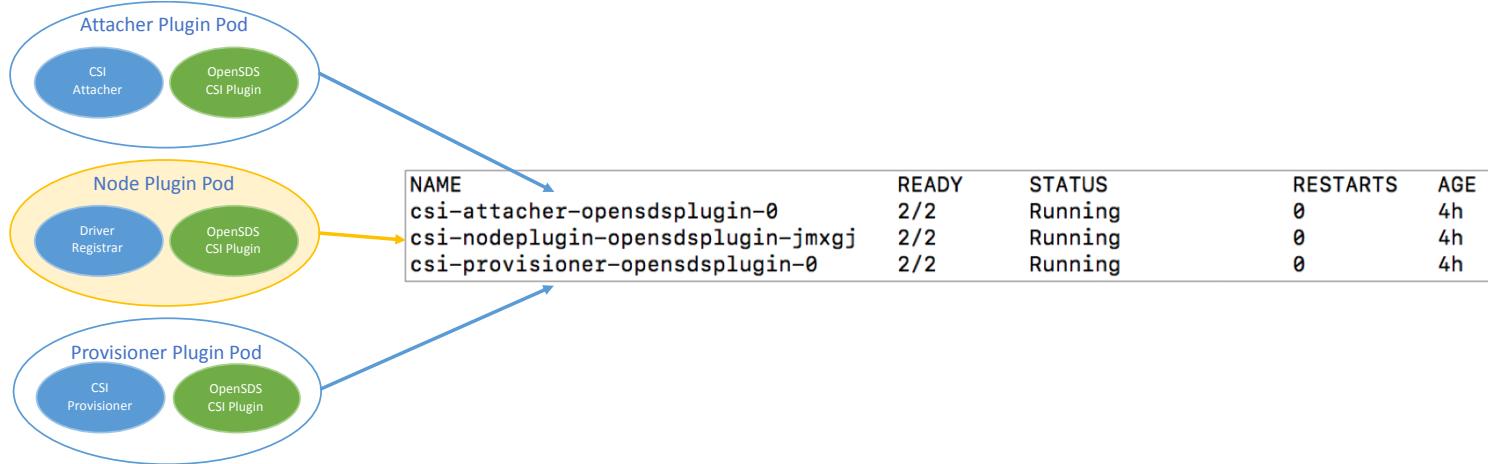
```
      storage: 10Gi
```

```
  storageClassName: opensds-csi-high-performance-sc
```

Note: profile parameter can be profile id or name

Running OpenSDS CSI Plugin

- Create OpenSDS CSI plugin pods:
`kubectl create -f csi/server/deploy/kubernetes`
- Three pods can be found by `kubectl get pod`:



Using OpenSDS Volume

- Create nginx application

```
kubectl create -f  
csi/server/examples/kubernetes/nginx.yaml
```

- An OpenSDS volume is mounted at `/var/lib/www/html`.

```
docker exec -it <nginx container id>  
/bin/bash
```

```
root@nginx:/# mount | grep html  
/dev/sda on /var/lib/www/html type ext4 (rw,relatime,data=ordered)
```

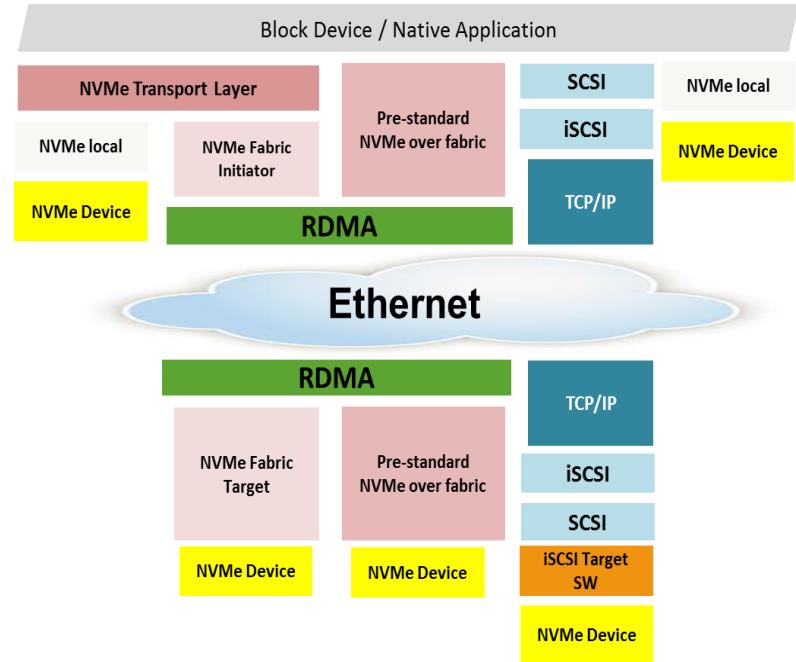
nginx.yaml

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: nginx  
spec:  
  containers:  
    - image: nginx  
      imagePullPolicy: IfNotPresent  
      name: nginx  
      ports:  
        - containerPort: 80  
          protocol: TCP  
      volumeMounts:  
        - mountPath: /var/lib/www/html  
          name: csi-data-opensdsplugin  
  volumes:  
    - name: csi-data-opensdsplugin  
      persistentVolumeClaim:  
        claimName: opensds-csi-high-performance-pvc  
        readOnly: false
```

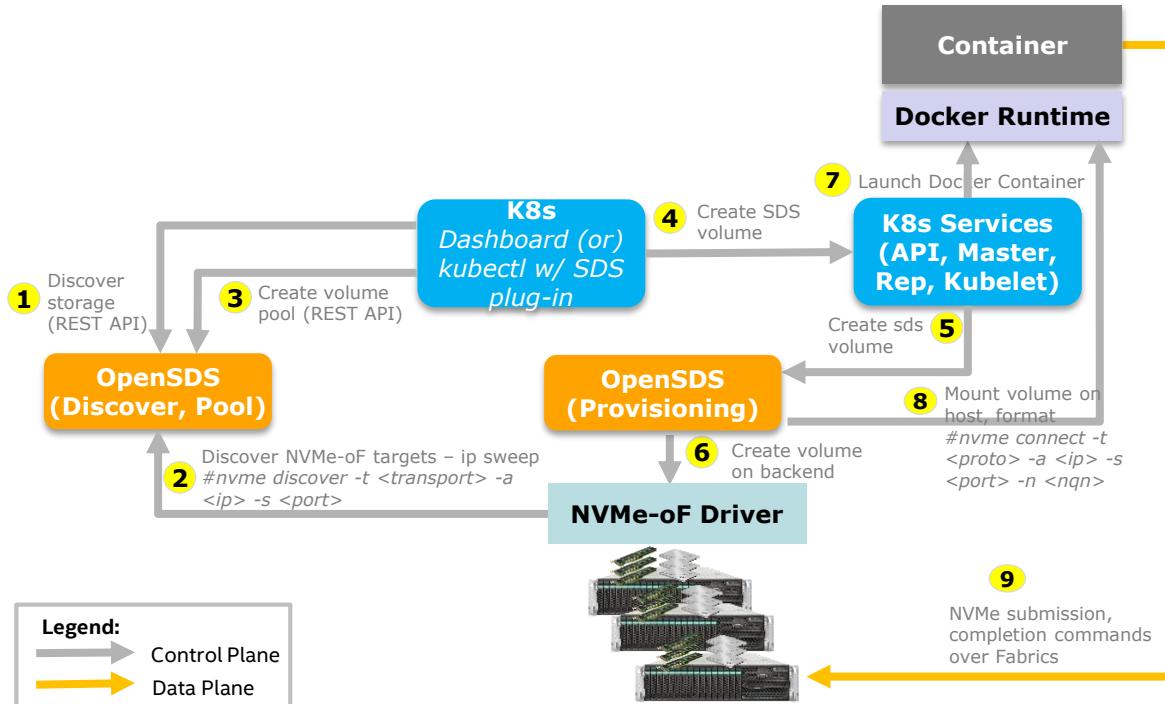
NVMe-oF: Local NVMe Performance

- The idea is to extend the efficiency of the local NVMe interface over a network fabric
 - Ethernet or IB
 - NVMe commands and data structures are transferred end to end
- Relies on RDMA for performance
 - Bypassing TCP/IP
- For more Information on NVMe over Fabrics (NVMe-oF)

[http://www.nvmeexpress.org/wp-content/uploads/NVMe Over Fabrics.pdf](http://www.nvmeexpress.org/wp-content/uploads/NVMe%20Over%20Fabrics.pdf)



NVMe-oF in OpenSDS



Work In Progress

- Linux Kernel Driver
- *nvme* commands for connect
- Drive assignment
- Profile includes target info

2019 Plans

- Multi-Target Pooling
- Rack aware scheduling
- Virtual Volume Provisioning
- NVMe over TCP/IP

OpenSDS Roadmap v0.18

2017H2 ZEALAND

- Kubernetes FlexVolume
- Vol CRUD
- Standalone Cinder Integration
- CSI Support
- Ceph, LVM

2018H1 ARUBA

- OpenStack
- Replication Array-Based, Host-Based
- Dashboard
- Storage Profiles
- Enumeration
- Block Storage
 - Cinder Drivers
 - Ceph
 - LVM
 - Huawei: Dorado

2018H2 BALI

- S3 Object
- Multi-Cloud Data Control
- Monitoring
- Storage Groups Snapshots, Replication
- NVMeoF Preview

2019H1 CAPRI*

- Multi-OpenStack
- File Share
- Migration
- Data Protection
- NVMeoF

2019H2++

- Optimization
- Tiering
- Security
- Sharing
- Networking

Governance

Technical Steering Committee



Steven Tan, Chair
Huawei, VP & CTO Cloud Storage Solution



Rakesh Jain, Vice-Chair
IBM, Research Engineer and Architect



Allen Samuels
Western Digital, R&D Engineering Fellow



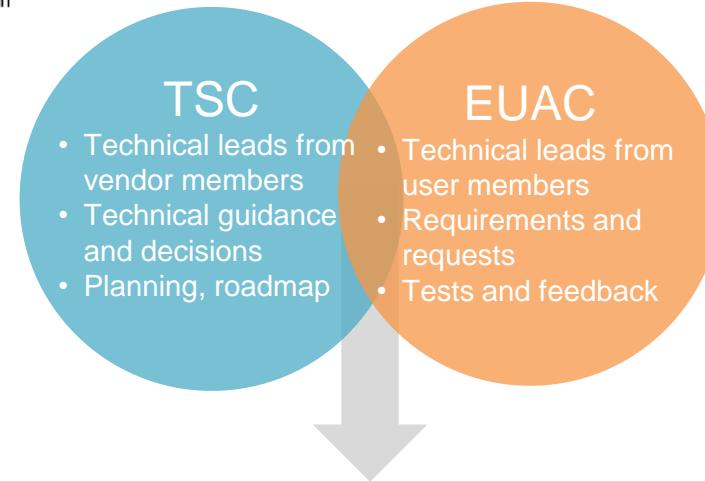
Anjaneya "Reddy" Chagam
Intel, Chief SDS Architect



Jay Bryant
Lenovo, Cloud Storage Lead



Shinya Hamano
Fujitsu, Software Architect



End-User Advisory Committee



Cosimo Rossetti
Vodafone, Lead Storage Architect



Yusuke Sato
Yahoo Japan, Infrastructure Lead



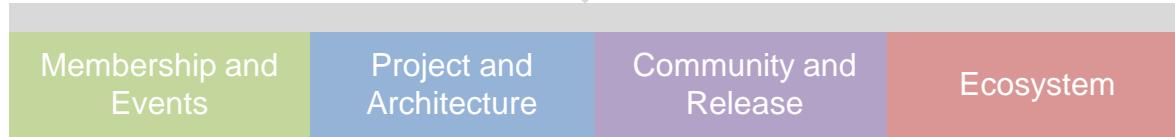
Kei Kusunoki
NTT Communications, Storage Architect



Yuji Yazawa
Toyota ITC, Group Lead



Wim Jacobs
KPN, Senior Architect



Summary

- ❑ OpenSDS aim is to address ***data storage integration challenges*** in cloud native environments with heterogeneous storage platforms
- ❑ OpenSDS is collaborating with SNIA to deliver Swordfish based DC manageability
- ❑ OpenSDS community is fast growing - join and contribute to address emerging storage management challenges

THANK YOU



-  <https://www.opensds.io>
-  <https://github.com/opensds>
-  info@opensds.io
-  [@opensds_io](https://twitter.com/opensds_io)
-  opensds.slack.com

FIND OUT
MORE

BE A
MEMBER

Meet us at the OpenSDS
booth!
Santa Clara, Sep 24-27

Accepting New Members
Vendors And End Users Welcome

Backup Slides

Join OpenSDS

End User Benefits

1. Nominate a representative to the End-User Advisory Committee (bi-weekly online EUAC meetings)
2. Leverage OpenSDS technologies to transform storage for cloud native era
3. Seek help and support from OpenSDS TSC, EUAC and developer community
4. Propose projects, specify requirements, and vote on roadmap priorities
5. Participate in OpenSDS sponsored events; speak and exhibit
6. Get vendors to work together



Join OpenSDS

Vendor Benefits

1. Nominate a representative to the Technical Steering Committee (bi-weekly online TSC meetings)
2. Engage end-users and understand their needs first hand
3. Propose projects, specify requirements, and guide the roadmap
4. Participate in OpenSDS sponsored events; speak and exhibit
5. Collaborate and explore potential partnerships with other vendors in an open manner

